

# Hadoop Tutorials Spark

Kacper Surdy  
Prasanth Kothuri

# About the tutorial


- The third session in Hadoop tutorial series
  - this time given by Kacper and Prasanth
- Session fully dedicated to Spark framework
  - Extensively discussed
  - Actively developed
  - Used in production
- Mixture of a talk and hands-on exercises

# What is Spark


- A framework for performing distributed computations
- Scalable, applicable for processing TBs of data
- Easy programming interface
- Supports Java, Scala, Python, R
- Varied APIs: DataFrames, SQL, MLlib, Streaming
- Multiple cluster deployment modes
- Multiple data sources HDFS, Cassandra, HBase, S3

# Compared to Impala

- Similar concept of the workload distribution
- Overlap in SQL functionalities
- Spark is more flexible and offers richer API
- Impala is fine tuned for SQL queries



# Evolution from MapReduce




# Deployment modes

## Local mode

- Make use of multiple cores/CPUs with the thread-level parallelism

## Cluster modes:

- Standalone
- Apache Mesos
- Hadoop YARN
  - typical for hadoop clusters



with centralised resource management


# How to use it

- Interactive shell
    - spark-shell
    - pyspark
  - Job submission
    - spark-submit (use also for python)
  - Notebooks
    - jupyter
    - zeppelin
    - SWAN (centralised CERN jupyter service)
- 
- terminal
- web interface

# Example – Pi estimation

```
import scala.math.random

val slices = 2
val n = 100000 * slices
val rdd = sc.parallelize(1 to n, slices)
val sample = rdd.map { i =>
    val x = random
    val y = random
    if (x*x + y*y < 1) 1 else 0
}
val count = sample.reduce(_ + _)
println("Pi is roughly " + 4.0 * count / n)
```



# Example – Pi estimation

- Login to one of the cluster nodes haperf1[01-12]
- Start spark-shell
- Copy or retype the content of  
`/afs/cern.ch/user/k/kasurdy/public/pi.spark`

# Spark concepts and terminology

# Transformations, actions (1)

- **Transformations** define how you want to transform your data. The result of a transformation of a spark dataset is another spark dataset. They do not trigger a computation.

examples: map, filter, union, distinct, join


- **Actions** are used to collect the result from a dataset defined previously. The result is usually an array or a number. They start the computation.

examples: reduce, collect, count, first, take

# Transformations, actions (2)

```
import scala.math.random


val slices = 2
val n = 100000 * slices
val rdd = sc.parallelize(1 to n, slices)
val sample = rdd.map { i =>
    val x = random
    val y = random
    if (x*x + y*y < 1) 1 else 0
}
val count = sample.reduce(_ + _)
println("Pi is roughly " + 4.0 * count / n)
```



# Driver, worker, executor

```
import scala.math.random


val slices = 2
val n = 100000 * slices
val rdd = sc.parallelize(1 to n, slices)
val sample = rdd.map { i =>
    val x = random
    val y = random
    if (x*x + y*y < 1) 1 else 0
}
val count = sample.reduce(_ + _)
println("Pi is roughly " + 4.0 * count / n)
```



# Stages, tasks

- Task - A unit of work that will be sent to one executor
- Stage - Each job gets divided into smaller sets of “tasks” called *stages* that depend on each other aka synchronization points

# Your job as a directed acyclic graph (DAG)



# Monitoring pages


- If you can, scroll the console output up to get an application ID (e.g. application\_1468968864481\_0070) and open

[http://haperf100.cern.ch:8088/proxy/application\\_XXXXXXXXXXXXXX\\_XXXX/](http://haperf100.cern.ch:8088/proxy/application_XXXXXXXXXXXXXX_XXXX/)

- Otherwise use the link below to find your application on the list

<http://haperf100.cern.ch:8088/cluster>

# Monitoring pages



# Data APIs